

gurobipy
Course

Xavier Nodet, xavier.nodet@gurobi.com

October 2025 - 0d492a7

Latest update of the course, slides and exercises



<https://nodet.github.io>

Gurobi can solve models with linear and quadratic (incl. nonconvex) constraints and objective.

$$\begin{array}{ll}\text{minimize} & x^T Q x + c^T x + d \\ \text{subject to} & Ax = b \\ & x^T Q_i x + c_i^T x \leq d_i \quad \forall i \in I \\ & l \leq x \leq u \\ & x_j \in \mathbb{Z} \quad \forall j \in J\end{array}$$

And nonlinear constraints as well!

1. Introductory examples

Let us start with a simple example

$$\begin{array}{ll}\text{maximize} & x + y + 2z \\ \text{subject to} & x + 2y + 3z \leq 4 \\ & x + y \geq 1 \\ & x, y, z \in \{0, 1\}\end{array}$$

Here is the Python code to solve this problem:

```
import gurobipy as gp ①
from gurobipy import GRB ②

with gp.Env() as env, gp.Model("simple-example", env=env) as model: ③
    x = model.addVar(vtype=GRB.BINARY, name="x") ④
    y = model.addVar(vtype=GRB.BINARY, name="y")
    z = model.addVar(vtype=GRB.BINARY, name="z")

    model.addConstr(x + 2 * y + 3 * z <= 4, name="c0") ⑤
    model.addConstr(x + y >= 1, name="c1")

    model.setObjective(x + y + 2 * z, sense=GRB.MAXIMIZE) ⑥

    model.write("example.lp") ⑦
```

```
model.optimize()
```

⑧

```
print("***** Solution *****")
```

```
for var in model.getVars():
```

```
    print(f"{var.VarName}: {var.X}")
```

⑨

```
print("*****")
```

- ① Import gurobipy package as gp for convenience
- ② GRB is the list of all Gurobi constants
- ③ Create a Gurobi environment and a model object
- ④ Define decision variables
- ⑤ Define constraints
- ⑥ Define objective

- ⑦ Save the model as an LP file
- ⑧ Optimize model
- ⑨ X attribute is the variable's value in the solution

Here is the log of the execution of this program:

```
Gurobi Optimizer version 12.0.3 build v12.0.3rc0 (mac64[arm] - Darwin 24.6.0
24G231)
```

```
CPU model: Apple M4 Pro
```

```
Thread count: 12 physical cores, 12 logical processors, using up to 12 threads
```

```
Optimize a model with 2 rows, 3 columns and 5 nonzeros
```

```
Model fingerprint: 0x98886187
```

```
Variable types: 0 continuous, 3 integer (3 binary)
```

```
Coefficient statistics:
```

```
Matrix range      [1e+00, 3e+00]
```

```
Objective range   [1e+00, 2e+00]
```

```
Bounds range      [1e+00, 1e+00]
```

```
RHS range         [1e+00, 4e+00]
```

Found heuristic solution: objective 2.0000000

Presolve removed 2 rows and 3 columns

Presolve time: 0.00s

Presolve: All rows and columns removed

Explored 0 nodes (0 simplex iterations) in 0.00 seconds (0.00 work units)

Thread count was 1 (of 12 available processors)

Solution count 2: 3 2

Optimal solution found (tolerance 1.00e-04)

Best objective 3.000000000000e+00, best bound 3.000000000000e+00, gap 0.0000%

***** Solution *****

x: 1.0

y: 0.0

z: 1.0

And here's the generated LP file:

```
\ Model simple-example
\ LP format - for model browsing. Use MPS format to capture full model detail.
Maximize
    x + y + 2 z
Subject To
    c0: x + 2 y + 3 z <= 4
    c1: x + y >= 1
Bounds
Binaries
    x y z
End
```

The same example, using the matrix API

```
import gurobipy as gp
from gurobipy import GRB
import numpy as np
import scipy.sparse as sp

with gp.Env() as env, gp.Model("matrix1", env=env) as m:

    # Create variables
    x = m.addMVar(shape=3, vtype=GRB.BINARY, name="x")

    # Set objective
    obj = np.array([1.0, 1.0, 2.0])
    m.setObjective(obj @ x, GRB.MAXIMIZE)
```

```
# Build (sparse) constraint matrix
row = np.array([0, 0, 0, 1, 1])
col = np.array([0, 1, 2, 0, 1])
val = np.array([1.0, 2.0, 3.0, -1.0, -1.0])
# A is such that A[row[k], col[k]] = val[k]
A = sp.csr_matrix((val, (row, col)), shape=(2, 3))

# Build rhs vector
rhs = np.array([4.0, -1.0])

# Add constraints
m.addConstr(A @ x <= rhs, name="c")

# Write the model
m.write("matrix1.lp")

# Optimize model
```

```
m.optimize()
```

```
print(x.X)
```

```
print(f"Obj: {m.ObjVal:g}")
```

Here's the log of the execution:

```
Gurobi Optimizer version 12.0.3 build v12.0.3rc0 (mac64[arm] - Darwin 24.6.0  
24G231)
```

```
CPU model: Apple M4 Pro
```

```
Thread count: 12 physical cores, 12 logical processors, using up to 12 threads
```

```
Optimize a model with 2 rows, 3 columns and 5 nonzeros
```

```
Model fingerprint: 0x8d4960d3
```

```
Variable types: 0 continuous, 3 integer (3 binary)
```

```
Coefficient statistics:
```

```
Matrix range      [1e+00, 3e+00]
```

```
Objective range   [1e+00, 2e+00]
```

```
Bounds range      [1e+00, 1e+00]
```

```
RHS range         [1e+00, 4e+00]
```


Found heuristic solution: objective 2.0000000

Presolve removed 2 rows and 3 columns

Presolve time: 0.00s

Presolve: All rows and columns removed

Explored 0 nodes (0 simplex iterations) in 0.00 seconds (0.00 work units)

Thread count was 1 (of 12 available processors)

Solution count 2: 3 2

Optimal solution found (tolerance 1.00e-04)

Best objective 3.000000000000e+00, best bound 3.000000000000e+00, gap 0.0000%

[1. 0. 1.]

Obj: 3

And here's the generated LP file:

```
\ Model matrix1
\ LP format - for model browsing. Use MPS format to capture full model detail.
Maximize
    x[0] + x[1] + 2 x[2]
Subject To
    c[0]: x[0] + 2 x[1] + 3 x[2] <= 4
    c[1]: - x[0] - x[1] <= -1
Bounds
Binaries
    x[0] x[1] x[2]
End
```

2. Python Data Structures

- **tuple**: An ordered, compound grouping that cannot be modified once it is created and is ideal for representing multi dimensional subscripts.

```
("city_0", "city_1")
```

- **list**: An ordered group, so each item is indexed. Lists can be modified by adding, deleting or sorting elements.

```
["city_0", "city_1", "city_2"]
```

- **set**: An unordered group of unique elements. Sets can only be modified by adding or deleting.

```
{"city_0", "city_1", "city_2"}
```

- **dict**: A key-value pair mapping that is ideal for representing indexed data such as cost, demand, capacity.

```
demand = {"city_0": 100, "city_1": 50, "city_2": 40}
```

3. Extended Data Structures in gurobipy

3.1. tuplelist

tuplelist: a sub-class of Python list, used to build sub-lists efficiently. See in particular **tuplelist.select(pattern)**.

```
import gurobipy as gp
l = gp.tuplelist([('A', 'B', 'C'),
                  ('A', 'C', 'D'),
                  ('A', 'E', 'C')])
print(l.select('A', '*', 'C'))
```

```
<gurobi.tuplelist (2 tuples, 3 values each):
( A , B , C )
```

(A , E , C)

>

3.2. tupledict

tupledict: a sub-class of Python dict, where the values are usually **Gurobi variables**, to efficiently retrieve those whose key match a specified tuple pattern.

Some important methods to build linear expressions efficiently:

- **tupledict.select(pattern)** → **list**
- **tupledict.sum(pattern)** → **gp.LinExpr**
- **tupledict.prod(coeff, pattern)** → **gp.LinExpr**

```
import gurobipy as gp
m = gp.Model()
```

```
x = m.addVars([(1,2), (1,3), (2,3)], name="x")
m.update()
expr = x.sum('*', 3)
print(expr)
```

```
# x is a tupledict
# Process all model updates
```

```
x[1,3] + x[2,3]
```


3.3. multidict()

`multidict()` is a convenience function to split a dict of lists.

```
import gurobipy as gp
keys, dict1, dict2 = gp.multidict( {
    'key1': [1, 2],
    'key2': [1, 3],
    'key3': [1, 4] } )
print(keys)
print(dict1)
print(dict2)
```

```
['key1', 'key2', 'key3']
{'key1': 1, 'key2': 1, 'key3': 1}
```

```
{'key1': 2, 'key2': 3, 'key3': 4}
```

3.4. Example of extended structures

```
import gurobipy as gp
from gurobipy import GRB

data = gp.tupledict([
    (("a", "b", "c"), 3),
    (("a", "c", "b"), 4),
    (("b", "a", "c"), 5),
    (("b", "c", "a"), 6),
    (("c", "a", "b"), 7),
    (("c", "b", "a"), 3)
])
print(f"data: {data}")

print("\nTuplelist:")
```

```

keys = gp.tuplelist(data.keys())
print(f"\tselect: {keys.select('a', '*', '*')}")

print("\nTupledict:")
print(f"\tselect  : {data.select('a', '*', '*')}")
print(f"\tsum      : {data.sum('*', '*', '*')}")
coeff = {("a", "c", "b"): 6, ("b", "c", "a"): -4}
print(f"\tprod     : {data.prod(coeff, '*', 'c', '*')}")

arcs, capacity, cost = gp.multidict({
    ("Detroit ", "Boston "): [100, 7],
    ("Detroit ", "New York "): [80, 5],
    ("Detroit ", "Seattle "): [120, 4],
    ("Denver ", "Boston "): [120, 8],
    ("Denver ", "New York "): [120, 11],
    ("Denver ", "Seattle "): [120, 4],
})

```

```
print("\nMultidict:")
print(f"\tcost: {cost}")
print("\n")
print(f"\tcapacity: {capacity}")
```

```
data: {('a', 'b', 'c'): 3, ('a', 'c', 'b'): 4, ('b', 'a', 'c'): 5, ('b', 'c', 'a'): 6, ('c', 'a', 'b'): 7, ('c', 'b', 'a'): 3}
```

Tuplelist:

```
select: <gurobi.tuplelist (2 tuples, 3 values each):
```

```
( a , b , c )
```

```
( a , c , b )
```

```
>
```

Tupledict:

```
select : [3, 4]
```

```
sum      : 28.0  
prod     : 0.0
```

Multidict:

```
cost: {('Detroit ', 'Boston '): 7, ('Detroit ', 'New York '): 5, ('Detroit  
, 'Seattle '): 4, ('Denver ', 'Boston '): 8, ('Denver ', 'New York '): 11,  
( 'Denver ', 'Seattle '): 4}
```

```
capacity: {('Detroit ', 'Boston '): 100, ('Detroit ', 'New York '): 80,  
( 'Detroit ', 'Seattle '): 120, ('Denver ', 'Boston '): 120, ('Denver ', 'New  
York '): 120, ('Denver ', 'Seattle '): 120}
```

4. Environments

Environments hold data that is global to one or more models.

- They hold a Gurobi license.
- They capture sets of parameter settings.
- They delineate a (single-threaded) [Gurobi session](#).

The basic usage pattern is the following:

```
import gurobipy as gp
from gurobipy import GRB

with gp.Env() as env, gp.Model("name", env=env) as m:
    # Use the model
    ...
```


A more advanced usage pattern is:

```
import gurobipy as gp
from gurobipy import GRB

with gp.Env(empty=True) as env:
    # Set licensing parameters
    env.setParam("CloudAccessID", "...")
    env.setParam("CloudSecretKey", "...")
    env.setParam("LicenseID", ...)
    # Start the environment before creating a model
    env.start()

    with gp.Model("name", env=env) as m:
        # Use the model
        ...
```

5. Models

A model holds:

- variables
- constraints
- parameters, that define the behavior of the solver

```
with gp.Env() as env, gp.Model("simple-example", env=env) as model:  
    x = model.addVar(vtype=GRB.BINARY, name="x")  
    y = model.addVar(vtype=GRB.BINARY, name="y")  
    c1 = model.addConstr(x + y >= 1, name="c1")  
    model.setObjective(x + y + 2 * z, sense=GRB.MAXIMIZE)  
  
    model.params.MipFocus=1  
    model.params.TimeLimit = 3600  
  
    model.optimize()
```

①

②

- ① Focus on finding the best possible solutions
- ② Stop after one hour

It has methods to create and edit variables and constraints, to set parameters, to solve the model, to retrieve information, and more.

```
# ...
```

```
model.write("model.mps")
```

①

```
model.chgCoeff(c1, x, 2)
```

②

```
model.computeIIS()
```

③

- ① Store the model in an MPS file
- ② Change the coefficient of variable x in constraint $c1$ to 2
- ③ Compute an Irreducible Inconsistent Set

5.1. Decision Variables, `Model.addVar()`

A decision variable is necessarily associated to exactly one instance of `Model`, and gets created using methods such as `Model.addVar()` to create a single variable, `Model.addVars()` to create multiple variables at once and `Model.addMVars()` to create a matrix of variables.

```
Model.addVar(lb=0.0, ub=float('inf'),  
             obj=0.0,  
             vtype=GRB.CONTINUOUS,  
             name="")
```

The available variable types in Gurobi are:

- Continuous: GRB.CONTINUOUS
- General integer: GRB.INTEGER
- Binary: GRB.BINARY
- Semi-continuous: GRB.SEMICONT
- Semi-integer: GRB.SEMIINT

A semi-continuous variable has the property that it takes a value of 0, or a value between the specified lower and upper bounds. A semi-integer variable adds the additional restriction that the variable should take an integral value.

```
# Define a binary decision variable with (default) lb=0
x = model.addVar(vtype=GRB.BINARY, name="x")
# Define an integer variable with lb=-1, ub=100
y = model.addVar(lb=-1, ub=100, vtype=GRB.INTEGER, name="y")
```

5.2. Model.addVars()

To add multiple decision variables to the model, use the `Model.addVars()` method which returns a Gurobi `tupledict` object containing the newly created variables:

```
Model.addVars(*indices,  
              lb=0.0, ub=float('inf'),  
              obj=0.0,  
              vtype=GRB.CONTINUOUS,  
              name="")
```


The first argument is an iterable giving indices for accessing the variables:

- several integers (specifying the dimensions of the matrix)
- several lists of scalars (each list specifies indices across one dimension of the matrix)
- one list of tuples, or a [tuplelist](#)

When the given name is a single string, it is subscripted by the index of the generator expression. The names are stored as ASCII strings, you should not use non-ASCII characters or spaces.

```
import gurobipy as gp
from gurobipy import GRB

with gp.Model(name="model") as model:
    # 3D array of binary variables
    x = model.addVars(2, 3, 4, vtype=GRB.BINARY, name="x")
    model.update()
    print(model.getAttr("VarName", model.getVars()))
```

```
['x[0,0,0]', 'x[0,0,1]', 'x[0,0,2]', 'x[0,0,3]', 'x[0,1,0]', 'x[0,1,1]',
 'x[0,1,2]', 'x[0,1,3]', 'x[0,2,0]', 'x[0,2,1]', 'x[0,2,2]', 'x[0,2,3]',
 'x[1,0,0]', 'x[1,0,1]', 'x[1,0,2]', 'x[1,0,3]', 'x[1,1,0]', 'x[1,1,1]',
 'x[1,1,2]', 'x[1,1,3]', 'x[1,2,0]', 'x[1,2,1]', 'x[1,2,2]', 'x[1,2,3]']
```

```
import gurobipy as gp
from gurobipy import GRB

with gp.Model(name="model") as model:
    # Use arbitrary lists of immutable objects -> tupledict
    y = model.addVars([1, 5], [7, 3, 2], ub=range(6),
                      name=[f"y_{i}" for i in range(6)])
    model.update()
    print("\nVariables names, upper bounds, and indices:")
    for index, var in y.items():
        print(f"name: {var.VarName}, ub: {var.UB}, index: {index}")
```

Variables names, upper bounds, and indices:

name: y_0, ub: 0.0, index: (1, 7)

name: y_1, ub: 1.0, index: (1, 3)

name: y_2, ub: 2.0, index: (1, 2)

name: y_3, ub: 3.0, index: (5, 7)

name: y_4, ub: 4.0, index: (5, 3)

name: y_5, ub: 5.0, index: (5, 2)

```
import gurobipy as gp
from gurobipy import GRB

with gp.Model(name="model") as model:
    # Use arbitrary list of tuples as indices
    z = model.addVars(
        [(3, "a"), (3, "b"), (7, "b"), (7, "c")], name="z",
    )
    model.update()
    print("\nVariables names and lower and upper bounds:")
    for index, var in z.items():
        print(f"name: {var.VarName}, lb: {var.LB}, ub: {var.UB}")
```

Variables names and lower and upper bounds:

name: z[3,a], lb: 0.0, ub: inf

name: z[3,b], lb: 0.0, ub: inf

name: z[7,b], lb: 0.0, ub: inf

name: z[7,c], lb: 0.0, ub: inf

5.3. Constraints, `Model.addConstr()`

Like variables, constraints are also associated with a model. Use the method `Model.addConstr()` to add a constraint to a model.

```
Model.addConstr(constr, name="")
```

`constr` is a `TempConstr` object that can take different types:

- Linear Constraint: $x + y \leq 1$
- Ranged Linear Constraint: $x + y == [1, 3]$
- Quadratic Constraint: $x*x + y*y + x*y \leq 1$
- Linear Matrix Constraint: $A @ x \leq 1$
- Quadratic Matrix Constraint: $x @ Q @ y \leq 2$
- Absolute Value Constraint: $x == \text{abs_}(y)$
- Logical Constraint: $x == \text{and_}(y, z)$
- Min or Max Constraint: $z == \text{max_}(x, y, \text{constant}=9)$
- Indicator Constraint: $(x == 1) \gg (y + z \leq 5)$


```

import gurobipy as gp
from gurobipy import GRB

# Add constraint "\sum_{i=0}^{n-1} x_i <= b" for any given n and b.
n, b = 10, 4
with gp.Model("model") as model:
    x = model.addVars(n, vtype=GRB.BINARY, name="x")
    c1 = model.addConstr(x.sum() <= b, name="c1")
    model.update()

    print(f"RHS, sense = {c1.RHS}, {c1.Sense}")
    print(f"row: {model.getRow(c1)}")

```

```

RHS, sense = 4.0, <
row: x[0] + x[1] + x[2] + x[3] + x[4] + x[5] + x[6] + x[7] + x[8] + x[9]

```

```

import gurobipy as gp
from gurobipy import GRB

# Add constraints "x_i + y_j - x_i * y_j >= 3".
n, m = 3, 2
with gp.Model("model") as model:
    x = model.addVars(n, name="x")
    y = model.addVars(m, name="y")
    for i in range(n):
        for j in range(m):
            model.addConstr(x[i] + y[j] - x[i] * y[j] >= 3, name=f"c_{i}{j}")
    model.update()

    for c in model.getQConstrs():
        print(f"Name: {c.QCName}, RHS: {c.QCRHS}, sense: {c.QCSense}")
        print(f"\trow: {model.getQCRow(c)}")

```

```
Name: c_00, RHS: 3.0, sense: >
    row: x[0] + y[0] + [ -1.0 x[0] * y[0] ]
Name: c_01, RHS: 3.0, sense: >
    row: x[0] + y[1] + [ -1.0 x[0] * y[1] ]
Name: c_10, RHS: 3.0, sense: >
    row: x[1] + y[0] + [ -1.0 x[1] * y[0] ]
Name: c_11, RHS: 3.0, sense: >
    row: x[1] + y[1] + [ -1.0 x[1] * y[1] ]
Name: c_20, RHS: 3.0, sense: >
    row: x[2] + y[0] + [ -1.0 x[2] * y[0] ]
Name: c_21, RHS: 3.0, sense: >
    row: x[2] + y[1] + [ -1.0 x[2] * y[1] ]
```

5.4. Model.addConstrs()

To add multiple constraints to the model, use the `Model.addConstrs()` method which returns a Gurobi `tupledict` that contains the newly created constraints:

```
Model.addConstrs(generator, name="")
```

```
import gurobipy as gp
from gurobipy import GRB

I = range(2)
J = ["a", "b", "c"]
with gp.Model("model") as model:
    x = model.addVars(I, name="x")
    y = model.addVars(J, name="y")

    # Add constraints  $x_i + y_j \leq 1$  for all (i, j)
    model.addConstrs((x[i] + y[j] <= 1 for i in I for j in J), name="c")
    model.update()
    print(model.getAttr("ConstrName", model.getConstrs()))
```

```
['c[0,a]', 'c[0,b]', 'c[0,c]', 'c[1,a]', 'c[1,b]', 'c[1,c]']
```

5.5. Objective Function

To set the model objective equal to a linear or a quadratic expression, use the `Model.setObjective()` method:

```
Model.setObjective(expr, sense=GRB.MINIMIZE)
```

`expr` can be:

- `LinExpr`, a linear expression
- `QuadExpr`, a quadratic expression

`sense` is either `GRB.MINIMIZE` (the default) or `GRB.MAXIMIZE`.

```
import gurobipy as gp
from gurobipy import GRB

import numpy as np

# Add linear objectives  $c^T x$ 
n = 5
c = np.random.rand(n)

with gp.Model("model") as model:
    x = model.addVars(n, name="x")
    linexpr = gp.quicksum(c_i * x_i for c_i, x_i in zip(c, x.values()))
    model.setObjective(linexpr)
    model.update()

    print(f"obj: {model.getObjective()}")
```

obj: 0.5262650100480577 x[0] + 0.31598755841732995 x[1] + 0.22589976835387182
x[2] + 0.26282494535866574 x[3] + 0.15402398277670748 x[4]


```
import gurobipy as gp
from gurobipy import GRB

import numpy as np

n = 5
Q = np.random.rand(n, n)

with gp.Model("model") as model:
    x = model.addVars(n, name="x")
    quadexpr = 0
    # Add quadratic objective in the form  $x^T Q x$ 
    for i in range(n):
        for j in range(n):
            quadexpr += x[i] * Q[i, j] * x[j]
    model.setObjective(quadexpr)
```

```
model.update()
```

```
# Print objective expression
```

```
obj = model.getObjective()
```

```
print(f"\nobj: {obj}")
```

```
obj: 0.0 + [ 0.35009519629328456 x[0] ^ 2 + 1.404423257261094 x[0] * x[1] +  
0.5633033648985756 x[0] * x[2] + 0.4734482637245059 x[0] * x[3] +  
0.9247814319759544 x[0] * x[4] + 0.3405485028613824 x[1] ^ 2 +  
0.44137792273617116 x[1] * x[2] + 0.6874279686129943 x[1] * x[3] +  
1.3001079227508634 x[1] * x[4] + 0.9103404225480711 x[2] ^ 2 +  
0.17321865502582912 x[2] * x[3] + 0.8862661425951331 x[2] * x[4] +  
0.3089399813193924 x[3] ^ 2 + 0.9932643804280612 x[3] * x[4] +  
0.7498300379876675 x[4] ^ 2 ]
```

5.6. Optimizing for Multiple Objectives

Gurobi supports two ways to **combine multiple linear objectives**:

- Blended objectives.
- Hierarchical objectives.

Objectives have:

- priorities
- weights
- absolute and relative tolerances

```
setObjectiveN(expr, index, priority=0, weight=1, abstol=1e-6, reltol=0, name='')
```

```
# Primary objective:  $x + 2y$ 
```

```
model.setObjectiveN(x + 2*y, 0, priority=0)
```

```
# Alternative, lower priority objectives:  $3y + z$  and  $x + z$ 
```

```
model.setObjectiveN(3*y + z, 1, priority=-1)
```

```
model.setObjectiveN(x + z, 2, priority=-2)
```

5.7. SOS Constraints

A Special-Ordered Set, or **SOS constraint**, is a highly specialized constraint that places restrictions on the values that variables in a given list can take.

- SOS constraint of type 1 (SOS1): at most one variable is allowed to take a non-zero value.
- SOS constraint of type 2 (SOS2): at most two variables are allowed to take non-zero values, and those non-zero variables must be contiguous.

Use `Model.addSOS()` to add such constraints:

```
Model.addSOS(type, vars)
```

With:

- *type*: the type of SOS constraint. Can be either `GRB.SOS_TYPE1` or `GRB.SOS_TYPE2`.
- *vars*: list of variables that participate in the constraint.

For example, the MIP formulation of

$$z = \max(x, y, 3)$$

using SOS1 constraints, is:

$$z = x + s_1 \quad (1)$$

$$z = y + s_2 \quad (2)$$

$$z = 3 + s_3 \quad (3)$$

$$v_1 + v_2 + v_3 = 1 \quad (4)$$

$$SOS1(s_1, v_1) \quad (5)$$

$$SOS1(s_2, v_2) \quad (6)$$

$$SOS1(s_3, v_3) \quad (7)$$

$$s_1, s_2, s_3 \in \mathbb{R}^+ \quad (8)$$

$$v_1, v_2, v_3 \in \{0, 1\} \quad (9)$$

5.8. General Constraints

General constraints allow you to directly model complex relationships between variables.

- Simple constraints: min, max, abs, OR, etc.

```
m.addConstr(z == gp.and_(x, y))  
m.addConstr(z == gp.max_(x, y, 3))
```

- Nonlinear constraints: polynomial, exponential, logistic, trigonometric, etc

```
model.addGenConstrNL(y, nlfunc.sin(2.5 * x1) + x2)
```

```
import gurobipy as gp
from gurobipy import nlfunc

# Minimize sin(2.5 x1) + x2
# s.t.      -1 <= x1, x2 <= 1

with gp.Env() as env, gp.Model(env=env) as model:
    x1 = model.addVar(lb=-1, ub=1, name="x1")
    x2 = model.addVar(lb=-1, ub=1, name="x2")

    y = model.addVar(lb=-float("inf"), name="y")
    model.addGenConstrNL(y, nlfunc.sin(2.5 * x1) + x2)
    model.setObjective(y)

    model.optimize()
    print(f"x1={x1.X}  x2={x2.X}  obj={y.X}")
```

Gurobi Optimizer version 12.0.3 build v12.0.3rc0 (mac64[arm] - Darwin 24.6.0 24G231)

CPU model: Apple M4 Pro

Thread count: 12 physical cores, 12 logical processors, using up to 12 threads

Optimize a model with 0 rows, 3 columns and 0 nonzeros

Model fingerprint: 0x00339c10

Model has 1 general nonlinear constraint (1 nonlinear terms)

Variable types: 3 continuous, 0 integer (0 binary)

Coefficient statistics:

Matrix range [0e+00, 0e+00]

Objective range [1e+00, 1e+00]

Bounds range [1e+00, 1e+00]

RHS range [0e+00, 0e+00]

Presolve model has 1 nlconstr

Added 2 variables to disaggregate expressions.

Presolve time: 0.00s

Presolved: 10 rows, 6 columns, 21 nonzeros

Presolved model has 1 nonlinear constraint(s)

Solving non-convex MINLP

Variable types: 6 continuous, 0 integer (0 binary)

Found heuristic solution: objective -2.0000000

Explored 1 nodes (0 simplex iterations) in 0.00 seconds (0.00 work units)

Thread count was 12 (of 12 available processors)

Solution count 1: -2

Optimal solution found (tolerance 1.00e-04)

Best objective -1.999999998349e+00, best bound -2.000000000000e+00, gap 0.0000%

$x_1 = -0.6282022274684884$ $x_2 = -1.0$ $obj = -1.9999999983490295$

6. Matrix-based API

- Term-based modeling can be slow
- Matrix-friendly API leans on Numpy concepts (vectorization, broadcasting)
- `Model.addMVar()`: Add an `MVar` object to a model. An `MVar` acts like a NumPy ndarray of Gurobi decision variables. An `MVar` can have an arbitrary number of dimensions, defined by the `shape` argument.

```
addMVar(shape, lb=0.0, ub=float('inf'), obj=0.0, vtype=GRB.CONTINUOUS, name='')
```

```
# A vector of 3 continuous variables
```

```
v = model.addMVar(3)
```

```
# A 5x10 matrix of binary variables
```

```
x = model.addMVar((5,10), vtype=GRB.BINARY)
```

Here's the example from the beginning of the course:

```
x = m.addMVar(shape=3, vtype=GRB.BINARY, name="x") ①
```

```
val = np.array([1.0, 2.0, 3.0, -1.0, -1.0])
```

```
row = np.array([0, 0, 0, 1, 1])
```

```
col = np.array([0, 1, 2, 0, 1])
```

```
A = sp.csr_matrix((val, (row, col)), shape=(2, 3)) ②
```

```
rhs = np.array([4.0, -1.0]) ③
```

```
m.addConstr(A @ x <= rhs, name="c") ④
```

```
obj = np.array([1.0, 1.0, 2.0])
```

```
m.setObjective(obj @ x, GRB.MAXIMIZE) ⑤
```

① Create variables

- ② Build a (sparse) matrix
- ③ Build an rhs vector
- ④ Add constraints
- ⑤ Set objective

Here is another with timing of the difference.

```
import gurobipy as gp
import numpy as np
from timeit import default_timer

n = 1000
Q = np.random.rand(n, n)

def term_based():
    with gp.Model("term-based") as model:
        x = model.addVars(n, name="x")
        model.addConstr(
            gp.quicksum(x[i] * Q[i, j] * x[j] for j in range(n)
                        for i in range(n)) <= 10
        )
```

```
def matrix_api():  
    with gp.Model("matrix-based") as model:  
        x = model.addMVar(n, name="x")  
        model.addConstr(x.T @ Q @ x <= 10)  
  
matrix_api() # To create the default env  
for f in [term_based, matrix_api]:  
    start = default_timer()  
    f()  
    end = default_timer()  
    print(f"Running {f.__name__} took {end - start} seconds")
```

```
Running term_based took 3.1411980830016546 seconds  
Running matrix_api took 0.073522040998796 seconds
```

7. Interacting with the Model

7.1. Attributes

The primary mechanism for querying and modifying properties of a Gurobi object is through the [attribute interface](#). Attributes exist on instances of Model, Variable, all types of constraints, and more.

Model:

- number of modeling elements of each type (**NumConstrs**, **NumVars**, etc.)
- information about the type of model (**IsMip**, **IsMultiObj**, etc.), its statistics (**SolCount**, **NodeCount**, etc.)
- information about the solutions found (**SolCount**), the best known bound (**ObjBound**), the gap (**MipGap**), etc.
- ...

Variables:

- lower (LB) and upper (UB) bounds
- value in a MIP start vector (Start)
- value in the best solution (X)
- ...

Constraints:

- right-hand side value (RHS)
- dual value in the best solution (Π)
- ...

```
import gurobipy as gp
from gurobipy import GRB

with gp.read("data/glass4.mps.bz2") as model:
    model.optimize()

    print("***** SOLUTION *****")
    print(f"\tStatus      : {model.Status}")
    print(f"\tObj          : {model.ObjVal}")
    print(f"\tSolutionCount: {model.SolCount}")
    print(f"\tRuntime       : {model.Runtime}")
    print(f"\tMIPGap        : {model.MIPGap}")

    print("\n")
    for var in model.getVars()[:20]:
        print(f"\t{var.VarName} = {var.X}")
```


Read MPS format model from file data/glass4.mps.bz2

Reading time = 0.03 seconds

glass4: 396 rows, 322 columns, 1815 nonzeros

Gurobi Optimizer version 12.0.3 build v12.0.3rc0 (mac64[arm] - Darwin 24.6.0 24G231)

CPU model: Apple M4 Pro

Thread count: 12 physical cores, 12 logical processors, using up to 12 threads

Optimize a model with 396 rows, 322 columns and 1815 nonzeros

Model fingerprint: 0x18b19fdf

Variable types: 20 continuous, 302 integer (0 binary)

Coefficient statistics:

Matrix range [1e+00, 8e+06]

Objective range [1e+00, 1e+06]

Bounds range [1e+00, 8e+02]

RHS range [1e+00, 8e+06]
 Presolve removed 6 rows and 6 columns
 Presolve time: 0.00s
 Presolved: 390 rows, 316 columns, 1803 nonzeros
 Variable types: 19 continuous, 297 integer (297 binary)
 Found heuristic solution: objective 3.133356e+09

Root relaxation: objective 8.000024e+08, 72 iterations, 0.00 seconds (0.00 work units)

Nodes			Current Node			Objective Bounds			Work	
Expl	Unexpl		Obj	Depth	IntInf	Incumbent	BestBd	Gap	It/Node	Time
	0	0	8.0000e+08	0	72	3.1334e+09	8.0000e+08	74.5%	-	0s
H	0	0				2.600019e+09	8.0000e+08	69.2%	-	0s
H	0	0				2.366684e+09	8.0000e+08	66.2%	-	0s
	0	0	8.0000e+08	0	72	2.3667e+09	8.0000e+08	66.2%	-	0s

	0	0	8.0000e+08	0	72	2.3667e+09	8.0000e+08	66.2%	-	0s
	0	0	8.0000e+08	0	77	2.3667e+09	8.0000e+08	66.2%	-	0s
	0	0	8.0000e+08	0	76	2.3667e+09	8.0000e+08	66.2%	-	0s
	0	2	8.0000e+08	0	75	2.3667e+09	8.0000e+08	66.2%	-	0s
H	27	53				2.000015e+09	8.0000e+08	60.0%	44.0	0s
H	28	53				2.000015e+09	8.0000e+08	60.0%	43.9	0s
H	132	173				1.991127e+09	8.0000e+08	59.8%	18.2	0s
H	275	285				1.977794e+09	8.0000e+08	59.6%	12.5	0s
H	1292	1340				1.966683e+09	8.0000e+08	59.3%	6.3	0s
H	1359	1338				1.900015e+09	8.0000e+08	57.9%	6.3	0s
H	2238	1917				1.900015e+09	8.0000e+08	57.9%	6.1	0s
H	2769	2163				1.800016e+09	8.0000e+08	55.6%	5.5	0s
H	3430	2418				1.800016e+09	8.0000e+08	55.6%	5.1	0s
H	3573	2282				1.750015e+09	8.0000e+08	54.3%	4.9	0s
*	5705	3167		75		1.740016e+09	8.0000e+08	54.0%	4.3	0s
*	5708	3075		76		1.722239e+09	8.0000e+08	53.5%	4.3	0s
H	5922	3078				1.700016e+09	8.0000e+08	52.9%	4.3	0s

* 6065	2959	67	1.700016e+09	8.0000e+08	52.9%	4.2	0s
* 7952	5105	93	1.700015e+09	8.0000e+08	52.9%	3.9	0s
H 9574	4973		1.666681e+09	8.0000e+08	52.0%	3.7	0s
H13664	7470		1.666680e+09	8.0000e+08	52.0%	3.6	0s
*16301	9211	79	1.600013e+09	8.0000e+08	50.0%	3.6	0s
H17437	9377		1.583346e+09	8.0000e+08	49.5%	3.6	0s
*19449	11033	54	1.575014e+09	8.0000e+08	49.2%	3.7	0s
*20221	10692	71	1.533348e+09	8.0000e+08	47.8%	3.8	0s
H22213	8731		1.475014e+09	8.0000e+08	45.8%	4.0	0s
H32855	12074		1.425014e+09	8.0001e+08	43.9%	3.7	0s
H35919	13342		1.400014e+09	9.0000e+08	35.7%	3.7	0s
*42393	13855	57	1.400013e+09	9.0401e+08	35.4%	5.3	2s
*73222	7138	71	1.375015e+09	1.1000e+09	20.0%	7.7	3s
*73524	3671	62	1.200013e+09	1.1000e+09	8.33%	7.8	3s

Cutting planes:

Learned: 1

Gomory: 7
Implied bound: 8
MIR: 26
Flow cover: 35
RLT: 4
Relax-and-lift: 17

Explored 76434 nodes (584415 simplex iterations) in 3.68 seconds (6.76 work units)

Thread count was 12 (of 12 available processors)

Solution count 10: 1.20001e+09 1.37501e+09 1.40001e+09 ... 1.60001e+09

Optimal solution found (tolerance 1.00e-04)

Best objective 1.200012600000e+09, best bound 1.200003400000e+09, gap 0.0008%

***** SOLUTION *****

Status : 2

Obj : 1200012600.0
SolutionCount: 10
Runtime : 3.6842780113220215
MIPGap : 7.666586172876269e-06

x1 = 0.0
x2 = 700.0
x3 = 1000.0
x4 = 1000.0
x5 = 400.0
x6 = 200.0
x7 = 200.0
x8 = 500.0
x9 = 700.0
x10 = 1200.0
x11 = 0.0

x12 = 0.0

x13 = 200.0

x14 = 800.0

x15 = 800.0

x16 = 600.0

x17 = 300.0

x18 = 300.0

x19 = 200.0

x20 = 1.0

7.2. Parameters

Parameters control the mechanics of the Gurobi Optimizer.

```
import gurobipy as gp
from gurobipy import GRB

with gp.read("data/glass4.mps.bz2") as model:
    model.params.Threads = 1
    model.params.TimeLimit = 10
    model.optimize()
```


Read MPS format model from file data/glass4.mps.bz2

Reading time = 0.03 seconds

glass4: 396 rows, 322 columns, 1815 nonzeros

Set parameter Threads to value 1

Set parameter TimeLimit to value 10

Gurobi Optimizer version 12.0.3 build v12.0.3rc0 (mac64[arm] - Darwin 24.6.0 24G231)

CPU model: Apple M4 Pro

Thread count: 12 physical cores, 12 logical processors, using up to 1 threads

Non-default parameters:

TimeLimit 10

Threads 1

Optimize a model with 396 rows, 322 columns and 1815 nonzeros

Model fingerprint: 0x18b19fdf

Variable types: 20 continuous, 302 integer (0 binary)

Coefficient statistics:

Matrix range [1e+00, 8e+06]

Objective range [1e+00, 1e+06]

Bounds range [1e+00, 8e+02]

RHS range [1e+00, 8e+06]

Presolve removed 6 rows and 6 columns

Presolve time: 0.00s

Presolved: 390 rows, 316 columns, 1803 nonzeros

Variable types: 19 continuous, 297 integer (297 binary)

Found heuristic solution: objective 3.133356e+09

Root relaxation: objective 8.000024e+08, 72 iterations, 0.00 seconds (0.00 work units)

Nodes		Current Node		Objective Bounds		Work
-------	--	--------------	--	------------------	--	------

Expl Unexpl			Obj	Depth	IntInf	Incumbent		BestBd	Gap	It/Node	Time
H	0	0	8.0000e+08	0	72	3.1334e+09	8.0000e+08	74.5%	-	0s	
	0	0				2.600019e+09	8.0000e+08	69.2%	-	0s	
	0	0	8.0000e+08	0	72	2.6000e+09	8.0000e+08	69.2%	-	0s	
	0	0	8.0000e+08	0	72	2.6000e+09	8.0000e+08	69.2%	-	0s	
	0	0	8.0000e+08	0	76	2.6000e+09	8.0000e+08	69.2%	-	0s	
H	0	0	8.0000e+08	0	76	2.6000e+09	8.0000e+08	69.2%	-	0s	
	0	0				2.500018e+09	8.0000e+08	68.0%	-	0s	
H	0	0				2.400019e+09	8.0000e+08	66.7%	-	0s	
	0	2	8.0000e+08	0	76	2.4000e+09	8.0000e+08	66.7%	-	0s	
H	52	52				2.288909e+09	8.0000e+08	65.0%	4.1	0s	
H	52	52				2.200018e+09	8.0000e+08	63.6%	4.1	0s	
H	52	52				2.150019e+09	8.0000e+08	62.8%	4.1	0s	
H	52	52				2.133352e+09	8.0000e+08	62.5%	4.1	0s	
H	54	54				2.000018e+09	8.0000e+08	60.0%	4.1	0s	
H	461	419				2.000018e+09	8.0000e+08	60.0%	4.1	0s	

H	461	419		2.000017e+09	8.0000e+08	60.0%	4.1	0s
H	461	411		1.950017e+09	8.0000e+08	59.0%	4.1	0s
H	461	411		1.933350e+09	8.0000e+08	58.6%	4.1	0s
H	461	411		1.900017e+09	8.0000e+08	57.9%	4.1	0s
H	547	479		1.900017e+09	8.0000e+08	57.9%	4.3	0s
H	688	536		1.900017e+09	8.0000e+08	57.9%	4.4	0s
H	708	526		1.900016e+09	8.0000e+08	57.9%	4.4	0s
H	708	501		1.866683e+09	8.0000e+08	57.1%	4.4	0s
*	735	483	126	1.833351e+09	8.0000e+08	56.4%	4.5	0s
H	795	495		1.833350e+09	8.0000e+08	56.4%	4.5	0s
H	795	475		1.800017e+09	8.0000e+08	55.6%	4.5	0s
H	795	457		1.800017e+09	8.0000e+08	55.6%	4.5	0s
H	958	509		1.800017e+09	8.0000e+08	55.6%	4.4	0s
*	1465	722	79	1.800015e+09	8.0000e+08	55.6%	4.3	0s
H	3508	2215		1.800015e+09	8.0000e+08	55.6%	3.9	0s
H	3508	2050		1.783348e+09	8.0000e+08	55.1%	3.9	0s
H	4508	2716		1.700014e+09	8.0000e+08	52.9%	3.9	0s

H 7156	4482				1.700014e+09	8.0000e+08	52.9%	4.0	1s
H 8150	5180				1.700014e+09	8.0000e+08	52.9%	3.9	1s
H10308	6336				1.700014e+09	8.4158e+08	50.5%	4.0	2s
H10308	6018				1.700014e+09	8.4158e+08	50.5%	4.0	2s
H10411	5783				1.700013e+09	8.5400e+08	49.8%	4.2	3s
H10411	5494				1.700013e+09	8.5400e+08	49.8%	4.2	3s
10530	5577	1.2000e+09	60	111	1.7000e+09	8.8337e+08	48.0%	4.8	5s
H10625	5358				1.650013e+09	8.9436e+08	45.8%	4.9	6s
H10631	5093				1.633347e+09	8.9439e+08	45.2%	4.9	6s
H10708	4888				1.600013e+09	8.9673e+08	44.0%	5.1	6s

Cutting planes:

Learned: 1

Gomory: 6

Implied bound: 7

Projected implied bound: 2

MIR: 35

Flow cover: 19

RLT: 3

Relax-and-lift: 12

Explored 17016 nodes (121739 simplex iterations) in 10.00 seconds (15.51 work units)

Thread count was 1 (of 12 available processors)

Solution count 10: 1.60001e+09 1.63335e+09 1.65001e+09 ... 1.70001e+09

Time limit reached

Best objective 1.600013400000e+09, best bound 9.000058582310e+08, gap 43.7501%

7.3. Callbacks

A callback is a user-defined function invoked by Gurobi while the optimization process is going on. They enable more control over the optimization. They can be used to:

- customize the termination of the solve
- add user cuts and lazy constraints
- add custom feasible solutions
- monitor the progress of optimization
- customize the optimization progress display

A callback must be a function (actually, any callable object) that accepts two arguments:

- `model`: the model that is being solved
- `where`: from where is the Gurobi Optimizer is the callback invoked (presolve, simplex, barrier, MIP, at a node, etc.)

A callback can query information from the solver using `Model.cbGet()`. The type of information that can be queried depends on from `where` the callback is invoked. For example, when `where == PRESOLVE`, you can query the number of rows removed using `what == PRE_ROWDEL`. All the callback code values

that you can query are listed [here](#).

```
import gurobipy as gp
from gurobipy import GRB
from functools import partial

# Used by the callback to store or retrieve information
class CallbackData:
    def __init__(self):
        self.invocations = 0

# The callback, the function that will be invoked
def mycallback(model, where, *, cbdata):
    # We're only interested in this case
    if where == GRB.Callback.MIP:
        cbdata.invocations += 1
```

```
# Get information from the model
nodecnt = model.cbGet(GRB.Callback.MIP_NODCNT)
if nodecnt > 100:
    # Terminate the search
    model.terminate()

# Create a model from an instance file
with gp.read("data/glass4.mps.bz2") as model:
    # Create the object to store/retrieve information
    cbdata = CallbackData()
    # Create a callable with two args: model and where
    callback_func = partial(mycallback, cbdata=cbdata)

    # Optimize, with the callback
    model.optimize(callback_func)
```

```
# Confirm the callback was invoked
```

```
print(f"MIP Callback was invoked {cbdata.invocations} times.")
```

Read MPS format model from file data/glass4.mps.bz2

Reading time = 0.02 seconds

glass4: 396 rows, 322 columns, 1815 nonzeros

Gurobi Optimizer version 12.0.3 build v12.0.3rc0 (mac64[arm] - Darwin 24.6.0 24G231)

CPU model: Apple M4 Pro

Thread count: 12 physical cores, 12 logical processors, using up to 12 threads

Optimize a model with 396 rows, 322 columns and 1815 nonzeros

Model fingerprint: 0x18b19fdf

Variable types: 20 continuous, 302 integer (0 binary)

Coefficient statistics:

Matrix range [1e+00, 8e+06]

Objective range [1e+00, 1e+06]

Bounds range [1e+00, 8e+02]

RHS range [1e+00, 8e+06]
 Presolve removed 6 rows and 6 columns
 Presolve time: 0.00s
 Presolved: 390 rows, 316 columns, 1803 nonzeros
 Variable types: 19 continuous, 297 integer (297 binary)
 Found heuristic solution: objective 3.133356e+09

Root relaxation: objective 8.000024e+08, 72 iterations, 0.00 seconds (0.00 work units)

Nodes			Current Node			Objective Bounds			Work	
Expl	Unexpl		Obj	Depth	IntInf	Incumbent	BestBd	Gap	It/Node	Time
	0	0	8.0000e+08	0	72	3.1334e+09	8.0000e+08	74.5%	-	0s
H	0	0				2.600019e+09	8.0000e+08	69.2%	-	0s
H	0	0				2.366684e+09	8.0000e+08	66.2%	-	0s
	0	0	8.0000e+08	0	72	2.3667e+09	8.0000e+08	66.2%	-	0s

	0	0	8.0000e+08	0	72	2.3667e+09	8.0000e+08	66.2%	-	0s
	0	0	8.0000e+08	0	77	2.3667e+09	8.0000e+08	66.2%	-	0s
	0	0	8.0000e+08	0	76	2.3667e+09	8.0000e+08	66.2%	-	0s
	0	2	8.0000e+08	0	75	2.3667e+09	8.0000e+08	66.2%	-	0s
H	27	53				2.000015e+09	8.0000e+08	60.0%	44.0	0s
H	28	53				2.000015e+09	8.0000e+08	60.0%	43.9	0s

Cutting planes:

Gomory: 91

Cover: 2

Implied bound: 40

MIR: 2

RLT: 72

Relax-and-lift: 44

PSD: 7

Explored 112 nodes (2811 simplex iterations) in 0.07 seconds (0.09 work units)

Thread count was 12 (of 12 available processors)

Solution count 5: 2.00002e+09 2.00002e+09 2.36668e+09 ... 3.13336e+09

Solve interrupted

Best objective 2.000015200000e+09, best bound 8.000035418612e+08, gap 60.0001%

User-callback calls 600, time in user-callback 0.00 sec

MIP Callback was invoked 117 times.

Other methods of interest are:

- `Model.cbGetNodeRel()`: retrieve the values of the variables in the node relaxation solution at the current node.
- `Model.cbGetSolution()`: retrieve the values of the variables in the new MIP solution.
- `Model.cbCut()`: add a new cutting plane to the model.
- `Model.cbLazy()`: add a new lazy constraint to the model.
- `Model.cbSetSolution()`: import a user-constructed solution into Gurobi.

8. Requirements and installation

8.1. Setup for Windows

8.1.1. Check for Anaconda 3.9 to 3.12

If you have Anaconda Python installed and this is version 3.9, 3.10, 3.11 or 3.12, just go ahead and run

```
conda install -c gurobi gurobipy
```

You can now go directly to section [Verification of the installation](#).

8.1.2. Check for the `py` launcher

Check whether you already have a 'core' Python version installed:

- open a Command Prompt and run

```
py --list
```

- If `py --list` tells you that you have version 3.9, 3.10, 3.11 or 3.12, you can go ahead to the [Installation of Gurobipy](#) section below.

- If you get

'py' is not recognized as an internal or external command,
operable program or batch file.

then you need to install Python 3.12, as per the next section.

8.1.3. Install Python 3.12

If you have none of 3.9, 3.10, 3.11 or 3.12, install the latest 3.12 Python version from <https://www.python.org/downloads/windows/>.

- Download the 'Windows installer (64-bit)', unless you have an ARM machine
- Run the installer (you don't need to do it as an Administrator)
- Make sure that 'Add Python 3.12 to the path' is NOT checked (so that your access to your existing Python version is not changed)

- open a Command Prompt and run

```
py --list
```

- Verify that you have 3.12 listed

For more details on the installation process, refer to <https://docs.python.org/3.12/using/windows.html>.

8.1.4. Create a venv and activate it

The reason to create a **venv** is to fix the Python version and isolate the work done in the venv to make sure it doesn't impact the rest of your system.

- In a Command Prompt, run

```
md gurobipy-course  
cd gurobipy-course  
py -3.12 -m venv venv --prompt "gurobipy-course"
```

Now that the venv has been created, you can activate it with one of the commands below. Note that you should always do this in

any new Command Prompt window that you start.

If you use PowerShell (e.g. your prompt starts with **PS C:...**), use

```
.\venv\Scripts\Activate.ps1
```

If you use **cmd** (no **PS** at the start of your prompt), use

```
venv\Scripts\activate
```

After activation of a virtual environment, **python** and **pip** *in that window* refer to that environment, not the default system-wide

installation. You can now continue in [Installation of Gurobipy](#).

8.2. Setup for macOS

I will assume here that you have already installed [Homebrew](#). Otherwise, follow the instructions there.

Install **pyenv** and Python 3.12

```
brew install pyenv  
pyenv install 3.12
```

Create a local directory for the course.

```
mkdir gurobipy-course  
cd gurobipy-course
```

Configure pyenv to use Python 3.12 when in this directory.

```
pyenv local 3.12
```

Create a virtual environment for the course and activate it. From now on, remember to always activate your venv in whatever shell you use.

```
python -m venv venv --prompt "gurobipy-course"  
source venv/bin/activate
```

I also highly recommend that you set the variable `PIP_REQUIRE_VIRTUALENV=true` in your environment. This forces `pip`

to fail when not running in a virtual environment.

8.3. Installation of Gurobipy

Confirm that your venv is active: you should see that your prompt is modified to mention 'gurobipy-course'. With your venv activated, confirm that

```
python --version
```

runs Python 3.12

You can now run

```
pip install gurobipy
```

8.4. Verification of the installation

You can check that gurobipy is correctly installed by running

```
python -c "import gurobipy as gp; m = gp.Model(); print(gp.GRB.VERSION_MAJOR)"
```

This should output one or a few lines about your license, then the Gurobi version number, which should be 12 at this time. An example could be:

```
Restricted license - for non-production use only - expires 2026-11-23  
12
```

8.5. Get a free Gurobi license

- Using your academic email address, create an account on <https://portal.gurobi.com>.
- **From a computer that belongs to your University**, log onto <https://portal.gurobi.com/iam/licenses/request?type=academic> and request a 'WLS Academic' license
- Once you see your license at <https://portal.gurobi.com/iam/licenses/list>, open <https://license.gurobi.com/manager/licenses/>, click 'Download' at the bottom of the license card, give a name and a description to the API key that you're about to create.

- When the 'API key created' dialog appears, click the 'Download' button and save the 'gurobi.lic' file.
- Copy that file to your laptop, in your user profile folder (typically `C:\Users\[name]`)
- Verify that running

```
python -c "import gurobipy as gp; m = gp.Model()"
```

does **not** give you any more the 'Restricted license' message.

For more details, refer to [this page](#).