

gurobipy
Course

Xavier Nodet, xavier.nodet@gurobi.com

January 2025 - bdc98cd

Latest update of the course, slides and exercises



<https://nodet.github.io>

Gurobi can solve models with linear and quadratic (incl. nonconvex) constraints and objective.

$$\begin{array}{ll}\text{minimize} & x^T Q x + c^T x + d \\ \text{subject to} & Ax = b \\ & x^T Q_i x + c_i^T x \leq d_i \quad \forall i \in I \\ & l \leq x \leq u \\ & x_j \in \mathbb{Z} \quad \forall j \in J\end{array}$$

And nonlinear constraints as well!

1. Introductory examples

Let us start with a simple example

$$\begin{array}{ll}\text{maximize} & x + y + 2z \\ \text{subject to} & x + 2y + 3z \leq 4 \\ & x + y \geq 1 \\ & x, y, z \in \{0, 1\}\end{array}$$

Here is the Python code to solve this problem:

```
import gurobipy as gp ①
from gurobipy import GRB ②

with gp.Env() as env, gp.Model("simple-example", env=env) as model: ③
    x = model.addVar(vtype=GRB.BINARY, name="x") ④
    y = model.addVar(vtype=GRB.BINARY, name="y")
    z = model.addVar(vtype=GRB.BINARY, name="z")

    model.addConstr(x + 2 * y + 3 * z <= 4, name="c0") ⑤
    model.addConstr(x + y >= 1, name="c1")

    model.setObjective(x + y + 2 * z, sense=GRB.MAXIMIZE) ⑥

    model.write("example.lp") ⑦
```

```
model.optimize()
```

⑧

```
print("***** Solution *****")
```

```
for var in model.getVars():
```

```
    print(f"{var.VarName}: {var.X}")
```

⑨

```
print("*****")
```

- ① Import gurobipy package as gp for convenience
- ② GRB is the list of all Gurobi constants
- ③ Create a Gurobi environment and a model object
- ④ Define decision variables
- ⑤ Define constraints
- ⑥ Define objective

- ⑦ Save the model as an LP file
- ⑧ Optimize model
- ⑨ X attribute is the variable's value in the solution

Here is the log of the execution of this program:

```
Gurobi Optimizer version 12.0.0 build v12.0.0rc1 (mac64[arm] - Darwin 24.2.0
24C101)
```

```
CPU model: Apple M1 Pro
```

```
Thread count: 8 physical cores, 8 logical processors, using up to 8 threads
```

```
Optimize a model with 2 rows, 3 columns and 5 nonzeros
```

```
Model fingerprint: 0x98886187
```

```
Variable types: 0 continuous, 3 integer (3 binary)
```

```
Coefficient statistics:
```

```
Matrix range      [1e+00, 3e+00]
```

```
Objective range   [1e+00, 2e+00]
```

```
Bounds range      [1e+00, 1e+00]
```

```
RHS range         [1e+00, 4e+00]
```

Found heuristic solution: objective 2.0000000

Presolve removed 2 rows and 3 columns

Presolve time: 0.00s

Presolve: All rows and columns removed

Explored 0 nodes (0 simplex iterations) in 0.00 seconds (0.00 work units)

Thread count was 1 (of 8 available processors)

Solution count 2: 3 2

Optimal solution found (tolerance 1.00e-04)

Best objective 3.000000000000e+00, best bound 3.000000000000e+00, gap 0.0000%

***** Solution *****

x: 1.0

y: 0.0

z: 1.0

And here's the generated LP file:

```
\ Model simple-example
\ LP format - for model browsing. Use MPS format to capture full model detail.
\ Signature: 0xd6af213f17f735ae
Maximize
    x + y + 2 z
Subject To
    c0: x + 2 y + 3 z <= 4
    c1: x + y >= 1
Bounds
Binaries
    x y z
End
```

The same example, using the matrix API

```
import gurobipy as gp
from gurobipy import GRB
import numpy as np
import scipy.sparse as sp

with gp.Env() as env, gp.Model("matrix1", env=env) as m:

    # Create variables
    x = m.addMVar(shape=3, vtype=GRB.BINARY, name="x")

    # Set objective
    obj = np.array([1.0, 1.0, 2.0])
    m.setObjective(obj @ x, GRB.MAXIMIZE)
```

```
# Build (sparse) constraint matrix
row = np.array([0, 0, 0, 1, 1])
col = np.array([0, 1, 2, 0, 1])
val = np.array([1.0, 2.0, 3.0, -1.0, -1.0])
# A is such that A[row[k], col[k]] = val[k]
A = sp.csr_matrix((val, (row, col)), shape=(2, 3))

# Build rhs vector
rhs = np.array([4.0, -1.0])

# Add constraints
m.addConstr(A @ x <= rhs, name="c")

# Write the model
m.write("matrix1.lp")

# Optimize model
```

```
m.optimize()
```

```
print(x.X)
```

```
print(f"Obj: {m.ObjVal:g}")
```

Here's the log of the execution:

```
Gurobi Optimizer version 12.0.0 build v12.0.0rc1 (mac64[arm] - Darwin 24.2.0  
24C101)
```

```
CPU model: Apple M1 Pro
```

```
Thread count: 8 physical cores, 8 logical processors, using up to 8 threads
```

```
Optimize a model with 2 rows, 3 columns and 5 nonzeros
```

```
Model fingerprint: 0x8d4960d3
```

```
Variable types: 0 continuous, 3 integer (3 binary)
```

```
Coefficient statistics:
```

```
Matrix range      [1e+00, 3e+00]
```

```
Objective range   [1e+00, 2e+00]
```

```
Bounds range      [1e+00, 1e+00]
```

```
RHS range         [1e+00, 4e+00]
```


Found heuristic solution: objective 2.0000000

Presolve removed 2 rows and 3 columns

Presolve time: 0.00s

Presolve: All rows and columns removed

Explored 0 nodes (0 simplex iterations) in 0.00 seconds (0.00 work units)

Thread count was 1 (of 8 available processors)

Solution count 2: 3 2

Optimal solution found (tolerance 1.00e-04)

Best objective 3.000000000000e+00, best bound 3.000000000000e+00, gap 0.0000%

[1. 0. 1.]

Obj: 3

And here's the generated LP file:

```
\ Model matrix1
\ LP format - for model browsing. Use MPS format to capture full model detail.
\ Signature: 0xd6af213f17f735ae
Maximize
    x[0] + x[1] + 2 x[2]
Subject To
    c[0]: x[0] + 2 x[1] + 3 x[2] <= 4
    c[1]: - x[0] - x[1] <= -1
Bounds
Binaries
    x[0] x[1] x[2]
End
```

2. Python Data Structures

- **tuple**: An ordered, compound grouping that cannot be modified once it is created and is ideal for representing multi dimensional subscripts.

```
("city_0", "city_1")
```

- **list**: An ordered group, so each item is indexed. Lists can be modified by adding, deleting or sorting elements.

```
["city_0", "city_1", "city_2"]
```

- **set**: An unordered group of unique elements. Sets can only be modified by adding or deleting.

```
{"city_0", "city_1", "city_2"}
```

- **dict**: A key-value pair mapping that is ideal for representing indexed data such as cost, demand, capacity.

```
demand = {"city_0": 100, "city_1": 50, "city_2": 40}
```

3. Extended Data Structures in gurobipy

3.1. tuplelist

tuplelist: a sub-class of Python list, used to build sub-lists efficiently. See in particular **tuplelist.select(pattern)**.

```
import gurobipy as gp
l = gp.tuplelist([('A', 'B', 'C'),
                  ('A', 'C', 'D'),
                  ('A', 'E', 'C')])
print(l.select('A', '*', 'C'))
```

```
<gurobi.tuplelist (2 tuples, 3 values each):
 ( A , B , C )
```

(A , E , C)

>

3.2. tupledict

tupledict: a sub-class of Python dict, where the values are usually **Gurobi variables**, to efficiently retrieve those whose key match a specified tuple pattern.

Some important methods to build linear expressions efficiently:

- **tupledict.select(pattern)** → **list**
- **tupledict.sum(pattern)** → **gp.LinExpr**
- **tupledict.prod(coeff, pattern)** → **gp.LinExpr**

```
import gurobipy as gp
m = gp.Model()
```

```
x = m.addVars([(1,2), (1,3), (2,3)], name="x")
m.update()
expr = x.sum('*', 3)
print(expr)
```

```
# x is a tupledict
# Process all model updates
```

```
x[1,3] + x[2,3]
```


3.3. multidict()

`multidict()` is a convenience function to split a dict of lists.

```
import gurobipy as gp
keys, dict1, dict2 = gp.multidict( {
    'key1': [1, 2],
    'key2': [1, 3],
    'key3': [1, 4] } )
print(keys)
print(dict1)
print(dict2)
```

```
['key1', 'key2', 'key3']
{'key1': 1, 'key2': 1, 'key3': 1}
```

```
{'key1': 2, 'key2': 3, 'key3': 4}
```

3.4. Example of extended structures

```
import gurobipy as gp
from gurobipy import GRB

data = gp.tupledict([
    (("a", "b", "c"), 3),
    (("a", "c", "b"), 4),
    (("b", "a", "c"), 5),
    (("b", "c", "a"), 6),
    (("c", "a", "b"), 7),
    (("c", "b", "a"), 3)
])
print(f"data: {data}")

print("\nTuplelist:")
```

```

keys = gp.tuplelist(data.keys())
print(f"\tselect: {keys.select('a', '*', '*')}")

print("\nTupledict:")
print(f"\tselect  : {data.select('a', '*', '*')}")
print(f"\tsum      : {data.sum('*', '*', '*')}")
coeff = {("a", "c", "b"): 6, ("b", "c", "a"): -4}
print(f"\tprod     : {data.prod(coeff, '*', 'c', '*')}")

arcs, capacity, cost = gp.multidict({
    ("Detroit ", "Boston "): [100, 7],
    ("Detroit ", "New York "): [80, 5],
    ("Detroit ", "Seattle "): [120, 4],
    ("Denver ", "Boston "): [120, 8],
    ("Denver ", "New York "): [120, 11],
    ("Denver ", "Seattle "): [120, 4],
    })

```

```
print("\nMultidict:")
print(f"\tcost: {cost}")
print("\n")
print(f"\tcapacity: {capacity}")
```

```
data: {('a', 'b', 'c'): 3, ('a', 'c', 'b'): 4, ('b', 'a', 'c'): 5, ('b', 'c', 'a'): 6, ('c', 'a', 'b'): 7, ('c', 'b', 'a'): 3}
```

Tuplelist:

```
select: <gurobi.tuplelist (2 tuples, 3 values each):
```

```
( a , b , c )
```

```
( a , c , b )
```

```
>
```

Tupledict:

```
select : [3, 4]
```

```
sum      : 28.0  
prod     : 0.0
```

Multidict:

```
cost: {('Detroit ', 'Boston '): 7, ('Detroit ', 'New York '): 5, ('Detroit  
, 'Seattle '): 4, ('Denver ', 'Boston '): 8, ('Denver ', 'New York '): 11,  
( 'Denver ', 'Seattle '): 4}
```

```
capacity: {('Detroit ', 'Boston '): 100, ('Detroit ', 'New York '): 80,  
( 'Detroit ', 'Seattle '): 120, ('Denver ', 'Boston '): 120, ('Denver ', 'New  
York '): 120, ('Denver ', 'Seattle '): 120}
```

4. Environments

Environments hold data that is global to one or more models.

- They hold a Gurobi license.
- They capture sets of parameter settings.
- They delineate a (single-threaded) [Gurobi session](#).

The basic usage pattern is the following:

```
import gurobipy as gp
from gurobipy import GRB

with gp.Env() as env, gp.Model("name", env=env) as m:
    # Use the model
    ...
```


A more advanced usage pattern is:

```
import gurobipy as gp
from gurobipy import GRB

with gp.Env(empty=True) as env:
    # Set licensing parameters
    env.setParam("CloudAccessID", "...")
    env.setParam("CloudSecretKey", "...")
    env.setParam("LicenseID", ...)
    # Start the environment before creating a model
    env.start()

    with gp.Model("name", env=env) as m:
        # Use the model
        ...
```

5. Models

A model holds:

- variables
- constraints
- parameters, that define the behavior of the solver

```
with gp.Env() as env, gp.Model("simple-example", env=env) as model:  
    x = model.addVar(vtype=GRB.BINARY, name="x")  
    y = model.addVar(vtype=GRB.BINARY, name="y")  
    c1 = model.addConstr(x + y >= 1, name="c1")  
    model.setObjective(x + y + 2 * z, sense=GRB.MAXIMIZE)  
  
    model.params.MipFocus=1  
    model.params.TimeLimit = 3600  
  
    model.optimize()
```

①

②

- ① Focus on finding the best possible solutions
- ② Stop after one hour

It has methods to create and edit variables and constraints, to set parameters, to solve the model, to retrieve information, and more.

```
# ...
```

```
model.write("model.mps")
```

①

```
model.chgCoeff(c1, x, 2)
```

②

```
model.computeIIS()
```

③

- ① Store the model in an MPS file
- ② Change the coefficient of variable x in constraint $c1$ to 2
- ③ Compute an Irreducible Inconsistent Set

5.1. Decision Variables, `Model.addVar()`

A decision variable is necessarily associated to exactly one instance of `Model`, and gets created using methods such as `Model.addVar()` to create a single variable, `Model.addVars()` to create multiple variables at once and `Model.addMVars()` to create a matrix of variables.

```
Model.addVar(lb=0.0, ub=float('inf'),  
             obj=0.0,  
             vtype=GRB.CONTINUOUS,  
             name="")
```

The available variable types in Gurobi are:

- Continuous: GRB.CONTINUOUS
- General integer: GRB.INTEGER
- Binary: GRB.BINARY
- Semi-continuous: GRB.SEMICONT
- Semi-integer: GRB.SEMIINT

A semi-continuous variable has the property that it takes a value of 0, or a value between the specified lower and upper bounds. A semi-integer variable adds the additional restriction that the variable should take an integral value.

```
# Define a binary decision variable with (default) lb=0
```

```
x = model.addVar(vtype=GRB.BINARY, name="x")  
# Define an integer variable with lb=-1, ub=100  
y = model.addVar(lb=-1, ub=100, vtype=GRB.INTEGER, name="y")
```

5.2. Model.addVars()

To add multiple decision variables to the model, use the `Model.addVars()` method which returns a Gurobi `tupledict` object containing the newly created variables:

```
Model.addVars(*indices,  
              lb=0.0, ub=float('inf'),  
              obj=0.0,  
              vtype=GRB.CONTINUOUS,
```

```
name="")
```

The first argument is an iterable giving indices for accessing the variables:

- several integers (specifying the dimensions of the matrix)
- several lists of scalars (each list specifies indices across one dimension of the matrix)
- one list of tuples, or a [tuplelist](#)

When the given name is a single string, it is subscripted by the index of the generator expression. The names are stored as ASCII

strings, you should not use non-ASCII characters or spaces.

```
import gurobipy as gp
from gurobipy import GRB

with gp.Model(name="model") as model:
    # 3D array of binary variables
    x = model.addVars(2, 3, 4, vtype=GRB.BINARY, name="x")
    model.update()
    print(model.getAttr("VarName", model.getVars()))

    # Use arbitrary lists of immutable objects to create a tupledict of 6
    variables
    y = model.addVars([1, 5], [7, 3, 2], ub=range(6), name=[f"y_{i}" for i in
range(6)])
    model.update()
```

```

print("\nVariables names, upper bounds, and indices:")
for index, var in y.items():
    print(f"name: {var.VarName}, ub: {var.UB}, index: {index}")

# Use arbitrary list of tuples as indices
z = model.addVars(
    [(3, "a"), (3, "b"), (7, "b"), (7, "c")],
    # Default lower bound is 0
    ub=GRB.INFINITY,
    name="z",
)
model.update()
print("\nVariables names and lower and upper bounds:")
for index, var in z.items():
    print(f"name: {var.VarName}, lb: {var.LB}, ub: {var.UB}")

```

```
['x[0,0,0]', 'x[0,0,1]', 'x[0,0,2]', 'x[0,0,3]', 'x[0,1,0]', 'x[0,1,1]',  
'x[0,1,2]', 'x[0,1,3]', 'x[0,2,0]', 'x[0,2,1]', 'x[0,2,2]', 'x[0,2,3]',  
'x[1,0,0]', 'x[1,0,1]', 'x[1,0,2]', 'x[1,0,3]', 'x[1,1,0]', 'x[1,1,1]',  
'x[1,1,2]', 'x[1,1,3]', 'x[1,2,0]', 'x[1,2,1]', 'x[1,2,2]', 'x[1,2,3]']
```

Variables names, upper bounds, and indices:

name: y_0, ub: 0.0, index: (1, 7)

name: y_1, ub: 1.0, index: (1, 3)

name: y_2, ub: 2.0, index: (1, 2)

name: y_3, ub: 3.0, index: (5, 7)

name: y_4, ub: 4.0, index: (5, 3)

name: y_5, ub: 5.0, index: (5, 2)

Variables names and lower and upper bounds:

name: z[3,a], lb: 0.0, ub: inf

name: z[3,b], lb: 0.0, ub: inf

```
name: z[7,b], lb: 0.0, ub: inf  
name: z[7,c], lb: 0.0, ub: inf
```

5.3. Constraints, `Model.addConstr()`

Like variables, constraints are also associated with a model. Use the method `Model.addConstr()` to add a constraint to a model.

```
Model.addConstr(constr, name="")
```

`constr` is a `TempConstr` object that can take different types:

- Linear Constraint: $x + y \leq 1$
- Ranged Linear Constraint: $x + y == [1, 3]$

- Quadratic Constraint: $x*x + y*y + x*y \leq 1$
- Linear Matrix Constraint: $A @ x \leq 1$
- Quadratic Matrix Constraint: $x @ Q @ y \leq 2$
- Absolute Value Constraint: $x == \text{abs_}(y)$
- Logical Constraint: $x == \text{and_}(y, z)$
- Min or Max Constraint: $z == \text{max_}(x, y, \text{constant}=9)$
- Indicator Constraint: $(x == 1) \gg (y + z \leq 5)$

```
import gurobipy as gp
from gurobipy import GRB
```

```
# Add constraint " $\sum_{i=0}^{n-1} x_i \leq b$ " for any given n and b.
```

```
n, b = 10, 4
```

```
with gp.Model("model") as model:
```

```
    x = model.addVars(n, vtype=GRB.BINARY, name="x")
```

```
    model.addConstr(x.sum() <= b, name="c1")
```

```
    model.update()
```

```
# Print the LHS, Sense, and RHS of c1
```

```
c1 = model.getConstrByName("c1")
```

```
print(f"RHS, sense = {c1.RHS}, {c1.Sense}")
```

```
print(f"row: {model.getRow(c1)}")
```

```
print("\n\n")
```

```
RHS, sense = 4.0, <
```

```
row: x[0] + x[1] + x[2] + x[3] + x[4] + x[5] + x[6] + x[7] + x[8] + x[9]
```

```

import gurobipy as gp
from gurobipy import GRB

# Add constraints "x_i + y_j - x_i * y_j >= 3".
n, m = 3, 2
with gp.Model("model") as model:
    x = model.addVars(n, name="x")
    y = model.addVars(m, name="y")
    for i in range(n):
        for j in range(m):
            model.addConstr(x[i] + y[j] - x[i] * y[j] >= 3, name=f"c_{i}{j}")
    model.update()

# Print the LHS, Sense, and RHS of all c_ij constraints
for c in model.getQConstrs():
    print(f"Name: {c.QCName}")

```

```
print(f"\tRHS, sense = {c.QCRHS}, {c.QCSense}")
print(f"\trow: {model.getQCRow(c)}")
```

Name: c_00

RHS, sense = 3.0, >

row: $x[0] + y[0] + [-1.0 \ x[0] * y[0]]$

Name: c_01

RHS, sense = 3.0, >

row: $x[0] + y[1] + [-1.0 \ x[0] * y[1]]$

Name: c_10

RHS, sense = 3.0, >

row: $x[1] + y[0] + [-1.0 \ x[1] * y[0]]$

Name: c_11

RHS, sense = 3.0, >

row: $x[1] + y[1] + [-1.0 \ x[1] * y[1]]$

Name: c_20


```
RHS, sense = 3.0, >
row: x[2] + y[0] + [ -1.0 x[2] * y[0] ]
Name: c_21
RHS, sense = 3.0, >
row: x[2] + y[1] + [ -1.0 x[2] * y[1] ]
```

5.4. Model.addConstrs()

To add multiple constraints to the model, use the `Model.addConstrs()` method which returns a Gurobi `tupledict` that contains the newly created constraints:

```
Model.addConstrs(generator, name="")
```

```
import gurobipy as gp
from gurobipy import GRB

# Add constraints  $x_i + y_j \leq 1$  for all  $(i, j)$ . Assume  $x_i$  and  $y_j$  are binary
# variables
I = range(5)
J = ["a", "b", "c"]
with gp.Model("model") as model:
    x = model.addVars(I, vtype=GRB.BINARY, name="x")
    y = model.addVars(J, vtype=GRB.BINARY, name="y")

    generator = (x[i] + y[j] <= 1 for i in I for j in J)
    model.addConstrs(generator, name="c")
    model.update()

# Print constraint names
```

```
print(model.getAttr("ConstrName", model.getConstrs()))
```

```
['c[0,a]', 'c[0,b]', 'c[0,c]', 'c[1,a]', 'c[1,b]', 'c[1,c]', 'c[2,a]', 'c[2,b]',  
'c[2,c]', 'c[3,a]', 'c[3,b]', 'c[3,c]', 'c[4,a]', 'c[4,b]', 'c[4,c]']
```

5.5. Objective Function

To set the model objective equal to a linear or a quadratic expression, use the [Model.setObjective\(\)](#) method:

```
Model.setObjective(expr, sense=GRB.MINIMIZE)
```

expr can be:

- `LinExpr`, a linear expression (constant plus pairs of coefficients/variables) - `QuadExpr`, a quadratic expression (linear expression plus triplets of coefficients/variables/variables)

`sense` is either `GRB.MINIMIZE` (the default) or `GRB.MAXIMIZE`.

```
import gurobipy as gp
from gurobipy import GRB

import numpy as np

# Add linear objectives  $c^T x$ 
n = 5
c = np.random.rand(n)
```

```
Q = np.random.rand(n, n)

with gp.Model("model") as model:
    x = model.addVars(n, name="x")
    linexpr = gp.quicksum(c_i * x_i for c_i, x_i in zip(c, x.values()))
    model.setObjective(linexpr)
    model.update()

    # Print objective expression
    obj = model.getObjective()
    print(f"obj: {obj}")
```

```
obj: 0.010290274887780781 x[0] + 0.6272652761106894 x[1] + 0.10292130911287034
x[2] + 0.6247962960605083 x[3] + 0.7858519376970178 x[4]
```

```
import gurobipy as gp
```

```
from gurobipy import GRB

import numpy as np

# Add quadratic objective in the form  $x^T Q x$ 

n = 5
c = np.random.rand(n)
Q = np.random.rand(n, n)

with gp.Model("model") as model:
    x = model.addVars(n, name="x")
    quadexpr = 0
    for i in range(n):
        for j in range(n):
            quadexpr += x[i] * Q[i, j] * x[j]
    model.setObjective(quadexpr)
```

```
model.update()
```

```
# Print objective expression
```

```
obj = model.getObjective()
```

```
print(f"\nobj: {obj}")
```

```
obj: 0.0 + [ 0.8698825927309295 x[0] ^ 2 + 0.9458957648341012 x[0] * x[1] +  
0.7166223092382553 x[0] * x[2] + 1.520489520868984 x[0] * x[3] +  
0.8674579249195624 x[0] * x[4] + 0.6099970447084597 x[1] ^ 2 +  
0.8427929095284881 x[1] * x[2] + 1.3333468758846756 x[1] * x[3] +  
0.5676182617835261 x[1] * x[4] + 0.3515305731557786 x[2] ^ 2 +  
0.5340370239545422 x[2] * x[3] + 0.48784490906406186 x[2] * x[4] +  
0.046807208146449675 x[3] ^ 2 + 0.8236951182404657 x[3] * x[4] +  
0.4959184970663434 x[4] ^ 2 ]
```

5.6. Optimizing for Multiple Objectives

Most practical business problems require compromises between multiple contradictory objectives. Gurobi supports two ways to **combine multiple linear objectives**:

- Blended objectives: the values of each objective is multiplied by a coefficient, and the sum of these weighted values is the value of the solution.
- Hierarchical objectives: Gurobi performs multiple passes, optimizing each time for one objective while preventing any degradation of the value of the objectives from the previous

passes.

You can build arbitrarily complex combinations of blended and hierarchical objectives: all the objectives that have the same priority are blended together, and the (groups of) objectives of different priority are handled in successive optimization passes.

Each objective also has a *weight* with which the value of the linear expression is multiplied. That weight defaults to 1, but can be changed to build combinations of blended objectives, or to change the sense of the optimization for a single objective.

Finally, objectives have tolerances (both absolute and relative).

These tolerances are used to relax the value of the objective in subsequent passes. For example, if the value of the first objective is 50 in the solution found during the first pass, and the relative tolerance for this objective is 0.1, the value of that objective in subsequent passes will be allowed to degrade by at most 5 units.

```
setObjectiveN(expr, index, priority=0, weight=1, abstol=1e-6, reltol=0, name='')
```

```
# Primary objective:  $x + 2y$ 
```

```
model.setObjectiveN(x + 2*y, 0, 0)
```

```
# Alternative, lower priority objectives:  $3y + z$  and  $x + z$ 
```

```
model.setObjectiveN(3*y + z, 1, -1)
```

```
model.setObjectiveN(x + z, 2, -2)
```

5.7. SOS Constraints

A Special-Ordered Set, or [SOS constraint](#), is a highly specialized constraint that places restrictions on the values that variables in a given list can take.

- SOS constraint of type 1 (SOS1): at most one variable is allowed to take a non-zero value.
- SOS constraint of type 2 (SOS2): at most two variables are allowed to take non-zero values, and those non-zero variables must be contiguous.

Use [Model.addSOS\(\)](#) to add such constraints:

```
Model.addSOS(type, vars)
```

With:

- *type*: the type of SOS constraint. Can be either `GRB.SOS_TYPE1` or `GRB.SOS_TYPE2`.
- *vars*: list of variables that participate in the constraint.

For example, the MIP formulation of $z = \max(x, y, 3)$, using SOS1 constraints, is:

$$z = x + s_1 \quad (1)$$

$$z = y + s_2 \quad (2)$$

$$z = 3 + s_3 \quad (3)$$

$$v_1 + v_2 + v_3 = 1 \quad (4)$$

$$SOS1(s_1, v_1) \quad (5)$$

$$SOS1(s_2, v_2) \quad (6)$$

$$SOS1(s_3, v_3) \quad (7)$$

$$s_1, s_2, s_3 > 0 \quad (8)$$

$$v_1, v_2, v_3 \in \{0, 1\} \quad (9)$$

5.8. General Constraints

General constraints allow you to directly model complex relationships between variables.

5.8.1. Simple Constraints

Simple constraints are a modeling convenience. They allow you to state fairly simple relationships between variables (min, max, absolute value, logical OR, etc.). Techniques for translating these constraints into lower-level modeling objects (typically using auxiliary binary variables and linear or SOS constraints) are well known and can be found in optimization modeling textbooks.

They include max, min, abs, and, or, norm, indicator and piecewise-linear constraints. In Python, you can use the [General Constraints Helper Functions](#) to create such constraints.

```
m.addConstr(z == gp.and_(x, y))  
m.addConstr(z == gp.max_(x, y, 3))
```

5.8.2. Nonlinear constraints

Nonlinear constraints are relationships of the form $y = f(x)$ where x is a vector of Gurobi variables, y is a single Gurobi variable and f is a function $f: \mathbb{R}^n \rightarrow \mathbb{R}$. The function may be composed of arithmetic operations (addition, subtraction, multiplication) and various

univariate nonlinear functions.

- Polynomial: $y = p_0x^n + p_1x^{n-1} + \dots + p_{n-1}x + p_n$
- Natural exponential: $y = e^x$
- Exponential: $y = a^x$, where $a > 0$ is the base for the exponential function.
- Natural logarithm: $y = \ln(x)$
- Logarithm: $y = \log_a(x)$ where $a > 0$ is the base for the logarithmic function.
- Logistic: $y = \frac{1}{1 + e^{-x}}$
- Power: $y = x^a$, where $x \geq 0$ for any fractional a and $x > 0$ for $a < 0$.

- Sine: $y = \sin(x)$
- Cosine: $y = \cos(x)$
- Tangent: $y = \tan(x)$

```
import gurobipy as gp
from gurobipy import nlfunc

# Formulate and solve the simple nonlinear model

# Minimize y
# s.t.      y = sin(2.5 x1) + x2
#          -1 <= x1, x2 <= 1

with gp.Env() as env, gp.Model(env=env) as model:
```

```
# Optimization variables
```

```
x1 = model.addVar(lb=-1, ub=1, name="x1")
```

```
x2 = model.addVar(lb=-1, ub=1, name="x2")
```

```
# Auxiliary resultant variable for general constraint
```

```
y = model.addVar(lb=-float("inf"), name="y")
```

```
# Nonlinear constraint for y
```

```
model.addGenConstrNL(y, nlfunc.sin(2.5 * x1) + x2)
```

```
# Use y for objective function
```

```
model.setObjective(y)
```

```
model.optimize()
```

```
print(f"x1={x1.X}  x2={x2.X}  obj={y.X}")
```

Gurobi Optimizer version 12.0.0 build v12.0.0rc1 (mac64[arm] - Darwin 24.2.0 24C101)

CPU model: Apple M1 Pro

Thread count: 8 physical cores, 8 logical processors, using up to 8 threads

Optimize a model with 0 rows, 3 columns and 0 nonzeros

Model fingerprint: 0x00339c10

Model has 1 general nonlinear constraint (1 nonlinear terms)

Variable types: 3 continuous, 0 integer (0 binary)

Coefficient statistics:

Matrix range [0e+00, 0e+00]

Objective range [1e+00, 1e+00]

Bounds range [1e+00, 1e+00]

RHS range [0e+00, 0e+00]

Presolve model has 1 nlconstr

Added 2 variables to disaggregate expressions.

Presolve time: 0.00s

Presolved: 10 rows, 6 columns, 21 nonzeros

Presolved model has 1 nonlinear constraint(s)

Solving non-convex MINLP

Variable types: 6 continuous, 0 integer (0 binary)

Found heuristic solution: objective -2.0000000

Explored 1 nodes (0 simplex iterations) in 0.00 seconds (0.00 work units)

Thread count was 8 (of 8 available processors)

Solution count 1: -2

Optimal solution found (tolerance 1.00e-04)

Best objective -1.999999998349e+00, best bound -2.000000000000e+00, gap 0.0000%

```
x1=-0.6282022274684884 x2=-1.0 obj=-1.9999999983490295
```

6. Matrix-based API

Term-based modeling where the variables and constraints are constructed one at a time can be time-consuming in Python.

Gurobi's matrix-friendly API enables matrix based operations which can be significantly faster than the term-based modeling. The Matrix API complements the capabilities of term-based modeling leaning on Numpy concepts and semantics such as vectorization and broadcasting.

The relevant objects/methods are:

- `Model.addMVar()`: Add an `MVar` object to a model. An `MVar` acts like a NumPy ndarray of Gurobi decision variables. An `MVar` can have an arbitrary number of dimensions, defined by the `shape` argument.

```
addMVar(shape, lb=0.0, ub=float('inf'), obj=0.0, vtype=GRB.CONTINUOUS, name='')
```

- Overloaded operators such as Python matrix multiply (`@`) build `MLinExpr` objects. Typically, you would multiply a 2-D matrix by a 1-D `MVar` object (e.g. `expr = A @ x`). Most arithmetic operators are supported on `MLinExpr` objects, including addition

and subtraction (e.g., $\text{expr} = A @ x - B @ y$), multiplication by a constant (e.g. $\text{expr} = 2 * A @ x$), and point-wise multiplication with an ndarray or a sparse matrix.

- Overloaded relational operators ($=$, $<=$ and $>=$) are used to build TempConstr objects from MLinExpr objects. For example, $A @ x <= 1$ and $A @ x == B @ y$ are both linear matrix constraints.
- Finally, `Model.addConstr()` is used to add that constraint to the model, possibly giving it a name.

```
addConstr(constr, name="")
```

An example of using the matrix-based API was given in [Matrix-API example](#), at the very beginning of this course.

Here is another with timing of the difference.

```
import gurobipy as gp
import numpy as np
from timeit import default_timer

# Build  $x^T Q x \leq 10$ 

n = 1000
Q = np.random.rand(n, n)

def term_based():
    with gp.Model("term-based") as model:
```



```

x = model.addVars(n, name="x")
model.addConstr(
    gp.quicksum(x[i] * Q[i, j] * x[j] for j in range(n)
                for i in range(n)) <= 10
)

```

```

def matrix_api():
    with gp.Model("matrix-based") as model:
        x = model.addMVar(n, name="x")
        model.addConstr(x.T @ Q @ x <= 10)

```

```

matrix_api() # To create the default env
for f in [term_based, matrix_api]:
    start = default_timer()
    f()
    end = default_timer()

```

```
print(f"Running {f.__name__} took {end - start} seconds")
```

```
Running term_based took 5.105839417024981 seconds  
Running matrix_api took 0.10690700000850484 seconds
```

7. Interacting with the Model

7.1. Attributes

The primary mechanism for querying and modifying properties of a Gurobi object is through the [attribute interface](#). Attributes exist on instances of Model, Variable, all types of constraints, and more. Here are some of the most commonly used:

Model:

- number of modelling elements of each type
- information about the type of model, its statistics
- information about the solutions found, the best known bound, the gap, etc.
- ...

Variables:

- lower and upper bounds
- value in a MIP start vector

- value in the best solution
- ...

Constraints:

- right-hand side value
- dual value in the best solution

```
import gurobipy as gp
from gurobipy import GRB

with gp.read("data/glass4.mps.bz2") as model:
    model.optimize()
```

```
print("***** SOLUTION *****")
print(f"\tStatus      : {model.Status}")
print(f"\tObj          : {model.ObjVal}")
print(f"\tSolutionCount: {model.SolCount}")
print(f"\tRuntime       : {model.Runtime}")
print(f"\tMIPGap        : {model.MIPGap}")

print("\n")
for var in model.getVars()[:20]:
    print(f"\t{var.VarName} = {var.X}")
```

```
Read MPS format model from file data/glass4.mps.bz2
Reading time = 0.01 seconds
glass4: 396 rows, 322 columns, 1815 nonzeros
Gurobi Optimizer version 12.0.0 build v12.0.0rc1 (mac64[arm] - Darwin 24.2.0
24C101)
```

CPU model: Apple M1 Pro

Thread count: 8 physical cores, 8 logical processors, using up to 8 threads

Optimize a model with 396 rows, 322 columns and 1815 nonzeros

Model fingerprint: 0x18b19fdf

Variable types: 20 continuous, 302 integer (0 binary)

Coefficient statistics:

Matrix range [1e+00, 8e+06]

Objective range [1e+00, 1e+06]

Bounds range [1e+00, 8e+02]

RHS range [1e+00, 8e+06]

Presolve removed 6 rows and 6 columns

Presolve time: 0.00s

Presolved: 390 rows, 316 columns, 1803 nonzeros

Variable types: 19 continuous, 297 integer (297 binary)

Found heuristic solution: objective 3.133356e+09

Root relaxation: objective 8.000024e+08, 72 iterations, 0.00 seconds (0.00 work units)

Nodes			Current Node			Objective Bounds			Work	
Expl	Unexpl		Obj	Depth	IntInf	Incumbent	BestBd	Gap	It/Node	Time
	0	0	8.0000e+08	0	72	3.1334e+09	8.0000e+08	74.5%	-	0s
H	0	0				2.600019e+09	8.0000e+08	69.2%	-	0s
H	0	0				2.366684e+09	8.0000e+08	66.2%	-	0s
	0	0	8.0000e+08	0	72	2.3667e+09	8.0000e+08	66.2%	-	0s
	0	0	8.0000e+08	0	72	2.3667e+09	8.0000e+08	66.2%	-	0s
	0	0	8.0000e+08	0	77	2.3667e+09	8.0000e+08	66.2%	-	0s
	0	0	8.0000e+08	0	76	2.3667e+09	8.0000e+08	66.2%	-	0s
	0	2	8.0000e+08	0	75	2.3667e+09	8.0000e+08	66.2%	-	0s
H	26	80				2.116683e+09	8.0000e+08	62.2%	38.4	0s
H	36	80				2.016683e+09	8.0000e+08	60.3%	30.6	0s

H	65	80		2.000015e+09	8.0000e+08	60.0%	20.1	0s
H	251	274		1.991127e+09	8.0000e+08	59.8%	10.8	0s
*	1128	1027	112	1.920014e+09	8.0000e+08	58.3%	6.3	0s
H	1695	1546		1.812517e+09	8.0000e+08	55.9%	6.2	0s
H	1788	1729		1.800017e+09	8.0000e+08	55.6%	6.1	0s
H	2010	1645		1.766684e+09	8.0000e+08	54.7%	6.0	0s
H	2353	1796		1.700017e+09	8.0000e+08	52.9%	5.6	0s
H	2399	1782		1.700017e+09	8.0000e+08	52.9%	5.6	0s
H	3109	1938		1.700016e+09	8.0000e+08	52.9%	5.4	0s
*	9626	5603	77	1.700016e+09	8.0000e+08	52.9%	3.9	0s
*	9627	5603	77	1.700016e+09	8.0000e+08	52.9%	3.9	0s
H10514	6175			1.700016e+09	8.0000e+08	52.9%	3.9	0s
H11625	7019			1.666683e+09	8.0000e+08	52.0%	3.8	0s
H11975	6974			1.650016e+09	8.0000e+08	51.5%	3.7	0s
H13781	7848			1.650016e+09	8.0000e+08	51.5%	3.6	0s
H13830	7848			1.650016e+09	8.0000e+08	51.5%	3.6	0s
H14108	8324			1.600015e+09	8.0000e+08	50.0%	3.6	0s

H16117	9426				1.600015e+09	8.0000e+08	50.0%	3.5	0s
H18938	10840				1.600014e+09	8.0000e+08	50.0%	3.5	1s
H19146	9570				1.500013e+09	8.0000e+08	46.7%	3.5	1s
H20592	10593				1.500013e+09	8.0000e+08	46.7%	3.5	1s
H30377	13890				1.500012e+09	8.3576e+08	44.3%	3.7	2s
30568	14024	1.1000e+09	66	89	1.5000e+09	8.8699e+08	40.9%	3.9	5s
30880	14242	1.0000e+09	27	114	1.5000e+09	8.9978e+08	40.0%	4.3	10s
39957	16368	1.0000e+09	177	54	1.5000e+09	9.2005e+08	38.7%	6.4	15s
*135241	22495		219		1.400016e+09	1.1000e+09	21.4%	7.0	19s
*136789	22885		210		1.400014e+09	1.1000e+09	21.4%	7.0	19s
*142250	22031		209		1.400013e+09	1.1098e+09	20.7%	7.0	19s
148470	22002	1.4000e+09	186	35	1.4000e+09	1.1750e+09	16.1%	7.0	20s
*186816	33080		211		1.400013e+09	1.2000e+09	14.3%	7.1	21s
*228748	44827		206		1.400013e+09	1.2000e+09	14.3%	6.9	23s
266386	52292	infeasible	203		1.4000e+09	1.2000e+09	14.3%	6.7	25s
*284522	5732		204		1.200013e+09	1.2000e+09	0.00%	6.6	25s

Cutting planes:

Learned: 1

Gomory: 7

Implied bound: 12

Projected implied bound: 1

MIR: 42

Flow cover: 17

RLT: 5

Relax-and-lift: 16

Explored 284537 nodes (1881113 simplex iterations) in 25.78 seconds (29.59 work units)

Thread count was 8 (of 8 available processors)

Solution count 10: 1.20001e+09 1.40001e+09 1.40001e+09 ... 1.60001e+09

Optimal solution found (tolerance 1.00e-04)

Best objective 1.200012600000e+09, best bound 1.200009038285e+09, gap 0.0003%

***** SOLUTION *****

Status : 2
Obj : 1200012600.0
SolutionCount: 10
Runtime : 25.77737784385681
MIPGap : 2.968064893289939e-06

x1 = 0.0
x2 = 700.00000000000001
x3 = 1000.0
x4 = 1000.0
x5 = 400.0
x6 = 200.0
x7 = 200.0
x8 = 500.0

```
x9 = 700.0  
x10 = 1200.0  
x11 = 0.0  
x12 = 0.0  
x13 = 200.0  
x14 = 800.0  
x15 = 800.0  
x16 = 600.0  
x17 = 300.0  
x18 = 300.0  
x19 = 199.99999999999999  
x20 = 1.0
```

7.2. Parameters

Parameters control the mechanics of the Gurobi Optimizer.

```
import gurobipy as gp
from gurobipy import GRB

with gp.read("data/glass4.mps.bz2") as model:
    model.params.Threads = 1
    model.params.TimeLimit = 10
    model.optimize()
```

```
Read MPS format model from file data/glass4.mps.bz2
Reading time = 0.02 seconds
glass4: 396 rows, 322 columns, 1815 nonzeros
Set parameter Threads to value 1
Set parameter TimeLimit to value 10
Gurobi Optimizer version 12.0.0 build v12.0.0rc1 (mac64[arm] - Darwin 24.2.0
24C101)
```

CPU model: Apple M1 Pro

Thread count: 8 physical cores, 8 logical processors, using up to 1 threads

Non-default parameters:

TimeLimit 10

Threads 1

Optimize a model with 396 rows, 322 columns and 1815 nonzeros

Model fingerprint: 0x18b19fdf

Variable types: 20 continuous, 302 integer (0 binary)

Coefficient statistics:

Matrix range [1e+00, 8e+06]

Objective range [1e+00, 1e+06]

Bounds range [1e+00, 8e+02]

RHS range [1e+00, 8e+06]

Presolve removed 6 rows and 6 columns

Presolve time: 0.00s

Presolved: 390 rows, 316 columns, 1803 nonzeros
 Variable types: 19 continuous, 297 integer (297 binary)
 Found heuristic solution: objective 3.133356e+09

Root relaxation: objective 8.000024e+08, 72 iterations, 0.00 seconds (0.00 work units)

Nodes			Current Node			Objective Bounds			Work	
Expl	Unexpl		Obj	Depth	IntInf	Incumbent	BestBd	Gap	It/Node	Time
H	0	0	8.0000e+08	0	72	3.1334e+09	8.0000e+08	74.5%	-	0s
	0	0				2.600019e+09	8.0000e+08	69.2%	-	0s
	0	0	8.0000e+08	0	72	2.6000e+09	8.0000e+08	69.2%	-	0s
	0	0	8.0000e+08	0	72	2.6000e+09	8.0000e+08	69.2%	-	0s
	0	0	8.0000e+08	0	76	2.6000e+09	8.0000e+08	69.2%	-	0s
	0	0	8.0000e+08	0	76	2.6000e+09	8.0000e+08	69.2%	-	0s
H	0	0				2.500018e+09	8.0000e+08	68.0%	-	0s

H	0	0			2.400019e+09	8.0000e+08	66.7%	-	0s	
	0	2	8.0000e+08	0	76	2.4000e+09	8.0000e+08	66.7%	-	0s
H	52	52			2.288909e+09	8.0000e+08	65.0%	4.1	0s	
H	52	52			2.200018e+09	8.0000e+08	63.6%	4.1	0s	
H	52	52			2.150019e+09	8.0000e+08	62.8%	4.1	0s	
H	52	52			2.133352e+09	8.0000e+08	62.5%	4.1	0s	
H	54	54			2.000018e+09	8.0000e+08	60.0%	4.1	0s	
H	461	419			2.000018e+09	8.0000e+08	60.0%	4.1	0s	
H	461	419			2.000017e+09	8.0000e+08	60.0%	4.1	0s	
H	461	411			1.950017e+09	8.0000e+08	59.0%	4.1	0s	
H	461	411			1.933350e+09	8.0000e+08	58.6%	4.1	0s	
H	461	411			1.900017e+09	8.0000e+08	57.9%	4.1	0s	
H	547	479			1.900017e+09	8.0000e+08	57.9%	4.3	0s	
H	688	536			1.900017e+09	8.0000e+08	57.9%	4.4	0s	
H	708	526			1.900016e+09	8.0000e+08	57.9%	4.4	0s	
H	708	501			1.866683e+09	8.0000e+08	57.1%	4.4	0s	
*	735	483		126	1.833351e+09	8.0000e+08	56.4%	4.5	0s	

H	931	559				1.833351e+09	8.0000e+08	56.4%	4.4	0s
H	931	538				1.800017e+09	8.0000e+08	55.6%	4.4	0s
H	931	520				1.800017e+09	8.0000e+08	55.6%	4.4	0s
H	958	518				1.800017e+09	8.0000e+08	55.6%	4.4	0s
H	958	502				1.800017e+09	8.0000e+08	55.6%	4.4	0s
*	1108	537		83		1.800015e+09	8.0000e+08	55.6%	4.6	0s
H	2106	1138				1.775015e+09	8.0000e+08	54.9%	4.1	0s
H	2555	1485				1.775015e+09	8.0000e+08	54.9%	4.0	0s
*	4189	2637		57		1.758351e+09	8.0000e+08	54.5%	3.9	1s
H	4665	2932				1.725017e+09	8.0000e+08	53.6%	3.9	1s
H	5665	3637				1.725017e+09	8.0000e+08	53.6%	3.9	1s
H	5665	3621				1.700016e+09	8.0000e+08	52.9%	3.9	1s
H	6130	3954				1.700016e+09	8.0000e+08	52.9%	3.9	1s
	10353	6940	1.2200e+09	56	123	1.7000e+09	8.4590e+08	50.2%	4.3	5s
H10418	6636					1.675016e+09	8.4985e+08	49.3%	4.4	5s
H10435	6314					1.650016e+09	8.5187e+08	48.4%	4.4	6s
H10435	5998					1.650016e+09	8.5187e+08	48.4%	4.4	6s

H10760	5883	1.600016e+09	8.8011e+08	45.0%	5.1	8s
H10801	5616	1.600016e+09	8.8011e+08	45.0%	5.2	8s
H10801	5339	1.600016e+09	8.8011e+08	45.0%	5.2	8s
H11173	5253	1.600016e+09	9.0000e+08	43.8%	5.7	9s

Cutting planes:

Learned: 1

Gomory: 11

Cover: 1

Implied bound: 5

Projected implied bound: 2

Clique: 1

MIR: 34

Flow cover: 20

RLT: 5

Relax-and-lift: 13

```
Explored 12753 nodes (79852 simplex iterations) in 10.00 seconds (11.07 work units)
```

```
Thread count was 1 (of 8 available processors)
```

```
Solution count 10: 1.60002e+09 1.60002e+09 1.60002e+09 ... 1.72502e+09
```

```
Time limit reached
```

```
Best objective 1.600015500000e+09, best bound 9.000054810120e+08, gap 43.7502%
```

7.3. Callbacks

A callback is a user-defined function invoked by Gurobi while the optimization process is going on. They enable more control over the optimization. They can be used to:

- customize the termination of the solve
- add user cuts and lazy constraints
- add custom feasible solutions
- monitor the progress of optimization
- customize the optimization progress display

A callback must be a function that accepts two arguments:

- **model**: the model that is being solved
- **where**: from where is the Gurobi Optimizer is the callback invoked (presolve, simplex, barrier, MIP, at a node, etc.)

A callback can query information from the solver using `Model.cbGet()`. The type of information that can be queried depends on from `where` the callback is invoked. For example, when `where == PRESOLVE`, you can query the number of rows removed using `what == PRE_ROWDEL`. All the callback code values that you can query are listed [here](#).

Other methods of interest are:

- `Model.cbGetNodeRel()`: retrieve the values of the variables in the node relaxation solution at the current node.
- `Model.cbGetSolution()`: retrieve the values of the variables in

the new MIP solution.

- `Model.cbCut()`: add a new cutting plane to the model.
- `Model.cbLazy()`: add a new lazy constraint to the model.
- `Model.cbSetSolution()`: import a user-constructed solution into Gurobi.